

Two-body low-energy scattering

Theory and implementation: every function, every constant, and why

Pedro Henrique Gesualdo Modesto

São Carlos Institute of Physics — University of São Paulo

August 26, 2026

Abstract

This document takes apart the two-body low-energy scattering code: every function, every term, every parameter, and the reason behind each choice. It is not a user manual — it is the explanation of *why*. Every numerical constant that appears in the code has here the calculation or the measurement that fixed it. The errors made and corrected are collected in Sec. 16, because they are the most instructive part.

Contents

1	What universality means here	3
2	Conventions and units	5
3	The one idea: do not integrate what is already known	6
4	The four potentials	6
4.1	The range R , in closed form	7
4.2	The units of the coefficients	7
4.3	A factor of two in Eq. (121)	7
5	Numerov: the integrator	8
6	The logarithmic grid	8
6.1	The problem	8
6.2	The calculation	9
7	The two starts	10
8	The slope at the edge	10
9	What the code measures: a and r_0	11
9.1	The scattering length	11
9.2	The normalisation, and why it is needed	12
9.3	The effective range	12
10	Counting bound states	13
11	The bound state	14
12	The inverse problem: tuning	14
12.1	Two nested loops	14
13	Every constant, one by one	15

14 The tabulated data	16
14.1 Scale invariance in <code>GUESS</code>	17
14.2 The constant that is tabulated nowhere	17
15 The checks	18
16 Autopsy: the errors made	18
17 Results	19
17.1 Binding energies	20
17.2 The helium Lennard-Jones, and what the wrong exponent cost	20
17.3 Accuracy achieved	21
18 The error budget	22
19 A potential with no free parameters	24
20 What is missing	25
A The three functions in full	26
A.1 <code>numerov</code> — the integrator	26
A.2 <code>scattering</code> — the three readings	27
A.3 <code>tune</code> — the inverse problem	27
B The data, with values	28

Revision history

date	what	why
2026-08-26	$C_6 \rightarrow L^4 C_6$, $C_{12} \rightarrow L^{10} C_{12}$	the old L^6 came from treating V as dimensionless. Fixed by L. Madeira. It also unblocked the helium Lennard-Jones potential which had been silently dropped (Sec. 17.2).
2026-08-26	scale-invariance test added	the property that was wrong is directly testable; it guards itself.
2026-08-26	9 orders of magnitude $\rightarrow$ 13	eV was being compared against K without converting to a common view point, Sec. 1.
2026-08-26	matching equation now carries its denominator	the notes showed $u'(R) + \kappa u(R)$; the code divided by $\max(u , u' , 1)$. Review point, Sec. 11.
2026-08-26	“the deuteron has a node” removed	the bound state is nodeless; the zero-energy solution is $u(r) = 0$ and carries a node. Review point, Sec. 11.
2026-08-26	spurious-node entry in the autopsy marked wrong	a zero product never satisfies < 0 , so that was never the cause. Review point, App. A2.
2026-08-26	“three smooth” $\rightarrow$ “three purely attractive”	the square well is discontinuous. Review point, Sec. 4.
2026-08-26	Pöschl-Teller: closed form for a only	r_0 and the zero-energy function are known only at unitarity. Review point, Sec. 4.
2026-08-26	Aziz HFD-B implemented, Sec. 19	first potential with no free parameters, so its (a, r_0) is fixed by condition, and it reproduces the published HFD-B data to four digits. Forced the unit conversion to be used and proposed that its core is soft. Review point, Sec. 18.
2026-08-26	CODATA constant for parameter-free potentials	the inferred $\hbar^2/2m_r$ is 0.20% low, which is 6% in a by $d \ln a / d \ln(\hbar^2/2m_r) = 48.7$ for a halo state. Sec. 19.
2026-08-26	unitarity given a dimensionless threshold	$R < a < 1e4$ replaced by $ r_0/a > \text{UNITARITY} = 10^{-5}$. The old bound hid a length. Works over three decades. Review point, Sec. 9.
2026-08-26	error budget added, Sec. 18	all six knobs from the review swept: <code>POINTS</code> , <code>V_CORE</code> , <code>R_START_FRACTION</code> , the edge stencil, <code>brentk</code> , and <code>units</code> . Review point, Sec. 18.
2026-08-26	<code>R_START_FRACTION</code> lifted to module level	it was inline in <code>numerov</code> ; a knob that cannot be turned cannot be audited.
2026-08-26	units written as definitions	$V_{\text{code}} = (m_r/\hbar^2)V_{\text{phys}}$ and the same for E , plus $[V] = \text{fm}^{-2}$ and the conversion for V . Five review points derived from this one. Sec. 2.
2026-08-26	finite range called a choice, not a property	only the square well has one; the others come from $ V > 10^{-12}$. Sec. 3.
2026-08-26	$[C_6] = \text{fm}^4$, $[C_{12}] = \text{fm}^{10}$ written	same fact as the scaling exponents, read off the units.
2026-08-26	r_0 no longer “the size of the potential”	it can be negative, and it is a statement about the wavefunction. Sec. 9.
2026-08-26	$x = \ln(r/r_*)$, $r_* = 1 \text{ fm}$	no logarithm of a dimensional quantity. App. A1.
2026-08-26	p_1, p_2 defined, with units	including that <code>LennardJones</code> takes C_6 first. App. B.
2026-08-26	every error labelled absolute or relative	plus r_0/a as the dimensionless measure, which explains why the zero range misses the deuteron by 36% and helium by 10%. Sec. 18.
2026-08-26	figure caption $10^{10} \rightarrow 10^5$	10^{10} is what the reference prescribes; the code uses 10^5 . Review point, Sec. 4.

1 What universality means here

At low energy, when the de Broglie wavelength is comparable to or larger than the range of the potential, the detailed shape of the interaction becomes invisible. Only two numbers survive — the **scattering length** a and the **effective range** r_0 — through the effective-range expansion

$$k \cot \delta_0(k) = -\frac{1}{a} + \frac{1}{2} r_0 k^2 + O(k^4). \quad (1)$$

This is what universality means: the same two-parameter description applies across physical domains that have nothing else in common. In the words of Ref. [3], “*the universal behavior in the low-energy limit enables us to draw a parallel between cold atomic gases and nuclear systems.*”

The two systems used throughout this work make the point concrete:

system	a	r_0	E (measured)	domain
^4He dimer	90.4 Å	8.0 Å	−1.62 mK	atomic
deuteron	5.4112 fm	1.7436 fm	−2.224 MeV	nuclear

Both quantities must be converted to common units before they are compared, which is the whole point of the exercise. With $k_B = 8.617 \times 10^{-5}$ eV/K, the deuteron sits at 2.224×10^6 eV and the helium dimer at 1.396×10^{-7} eV: a ratio of 1.59×10^{13} , so **13 orders of magnitude in energy**. In length, 90.4 Å against 5.4112 fm is 1.67×10^6 , so **6 orders**. One is held together by van der Waals and the other by the strong interaction, and yet the same two numbers, fed into the same two formulas, reproduce both binding energies.

The two figures are not independent. If $E \sim \hbar^2/(2m_r a^2)$, the energy ratio should be the length ratio squared times the reduced-mass ratio: $(1.67 \times 10^6)^2 \times 3.97 = 1.11 \times 10^{13}$, against 1.59×10^{13} measured. The factor of 1.44 left over is what the effective range adds on top of the zero-range estimate. So the 13 orders in energy *are* the 6 in length, squared, plus the masses — which is the scaling law working, not a coincidence.

Shape independence is the mechanism, not the definition

A separate — and often confused — statement is that *different potentials* tuned to the same (a, r_0) give the same low-energy observables. That is not the definition of universality; it is a consequence of the **shape-independent approximation**, Eq. (55) of [3], which says that $\delta_0(k)$ at low energy does not depend on the shape of the potential. As [3] puts it, the fact that four very different potentials can be tuned to the same a and r_0 , “*although remarkable, is directly related to the shape-independent approximation.*”

So the roles are:

universality	two parameters describe low-energy physics across domains
shape independence	why the shape of V beyond those two numbers does not matter
tuning	the numerical procedure that puts a given V at a chosen (a, r_0)

Why: Tuning is a *demonstration technique*, not a physical claim. The tuned parameters are of course different for each potential — a Gaussian and a Lennard-Jones that share (a, r_0) have nothing numerically in common. What they share is everything a low-energy experiment can measure.

When two parameters are not enough

The expansion (1) is a series, and dropping r_0 means keeping only its first term. Whether that is legitimate depends on how small kR is. Reference [3] makes exactly this point for the deuteron: the zero-range formula “*only accounts for $\approx 64\%$ of the deuteron energy*”, because “ *$kR \sim 0.4$ for the deuteron is larger than in the ^4He case.*”

system	$r_0/ a $	zero range	finite range
deuteron	0.32	−36%	−0.05%
^4He dimer	0.089	−8.6%	+0.47%

Why: The deuteron is the textbook example of a shallow bound state, and it is the one that needs the finite-range correction most. Being universal is not a property a system has or lacks in the abstract; it is a statement about how small the expansion parameter happens to be.

2 Conventions and units

The s -wave radial Schrödinger equation, written for $u(r) = r \psi(r)$, is

$$-\frac{\hbar^2}{2m_r} u''(r) + V(r) u(r) = E u(r), \quad (2)$$

with $m_r = m_1 m_2 / (m_1 + m_2)$ the reduced mass. The code works with rescaled quantities, defined explicitly as

$$V_{\text{code}} = \frac{m_r}{\hbar^2} V_{\text{phys}}, \quad E_{\text{code}} = \frac{m_r}{\hbar^2} E_{\text{phys}}, \quad (3)$$

which turns the equation into

$$\boxed{u'' = 2(V - E)u} \quad (4)$$

the only equation the code ever solves, at $E = 0$ and at $E < 0$ alike.

Why: It is tempting to summarise this as “ $\hbar = m_r = 1$, lengths in fm”. That is imprecise, and the imprecision has consequences. Distances stay in fm; they are not rescaled. So Eq. (3) does *not* make V_{code} and E_{code} dimensionless — it gives them units of inverse length squared:

$$[V_{\text{code}}] = [E_{\text{code}}] = \text{fm}^{-2}.$$

The honest one-line statement is: *distances are expressed in fm, while energies and potentials are rescaled by $m_r/\hbar^2$, so their numerical units are fm^{-2}* . Everything about the Lennard-Jones coefficients — their units in Sec. 4, and the exponents in the scaling rule of Sec. 14.1 — follows from that one line. Treating V as dimensionless is exactly how the wrong exponent $C_{12} \rightarrow C_{12}L^6$ got into an earlier draft.

Back to physical units

Comparing (4) with the dimensional form,

$$\frac{2m_r}{\hbar^2} (V_{\text{phys}} - E_{\text{phys}}) = 2(V_{\text{code}} - E_{\text{code}}) \implies E_{\text{phys}} = \frac{\hbar^2}{m_r} E_{\text{code}},$$

and since $\hbar^2/2m_r$ is the constant usually tabulated,

$$\boxed{E_{\text{phys}} = 2 \left(\frac{\hbar^2}{2m_r} \right) E_{\text{code}}} \quad (5)$$

Why: That factor of 2 is easy to lose. The check that would expose it is in Sec. 14: with it wrong, the inferred constant would not come out at 41.47 MeV·fm².

The potential converts by the same rule, and it is worth writing separately because it is what the Lennard-Jones discussion of Sec. 4 turns on:

$$V_{\text{phys}}(r) = \frac{\hbar^2}{m_r} V_{\text{code}}(r) = 2 \left(\frac{\hbar^2}{2m_r} \right) V_{\text{code}}(r). \quad (6)$$

3 The one idea: do not integrate what is already known

Every potential here has a range R . Beyond R we have $V = 0$, and there Eq. (4) has a solution known without computing anything:

$$E = 0 : \quad u'' = 0 \quad \implies \quad u \text{ is a **straight line**,} \quad (7)$$

$$E < 0 : \quad u'' = -2Eu \quad \implies \quad u \propto e^{-\kappa r}, \quad \kappa = \sqrt{-2E}. \quad (8)$$

So the code **never integrates past R** . It integrates from the inner boundary out to R , reads the value and the slope there, and glues the analytic form on. All three quantities come out of that gluing:

quantity	comes from
a	where the asymptotic straight line crosses zero
r_0	the integral of the difference between that line and the true solution
E	the energy that makes the solution glue onto the exponential

Why: Integrating beyond R would mean computing numerically something already known exactly, and paying for it in numerical error.

4 The four potentials

The choice is not decorative: each shape stresses a different numerical hypothesis.

Potential	Form	What it stresses
Square well, Eq. (70)	$-v\mu^2$ for $r < 1/\mu$	a discontinuity
Pöschl-Teller, Eq. (116)	$-v\mu^2 / \cosh^2(\mu r)$	smooth ; closed form for a only
Gaussian, Eq. (120)	$-v\mu^2 e^{-\mu^2 r^2}$	tail dying faster than exponentially
Lennard-Jones, Eq. (121)	$\frac{1}{2}(C_{12}/r^{12} - C_6/r^6)$	hard core plus van der Waals tail




Figure 1: The four potentials, all tuned to unitarity ($|a| \rightarrow \infty$). The Lennard-Jones needs logarithmic axes because of its hard core; the dotted line marks $r_{\min}$, the radius where V reaches $V_{\text{CORE}} = 10^5$ and integration starts. Reference [3] prescribes 10^{10} for that cut; Sec. 13 shows why 10^5 is already converged.

4.1 The range R , in closed form

Each class needs to know where its potential stops acting. **Only one of the four actually has such a radius.** The square well is the exception: it is identically zero beyond $R = 1/\mu$, so its range is a property of the potential. The other three never truly vanish — a Gaussian decays forever, a van der Waals tail falls off as a power law — and for them a finite range does not exist as a mathematical fact. What exists is a *choice*:

$$|V(R)| = V_ZERO = 10^{-12}, \quad (9)$$

a truncation we impose because the integration has to stop somewhere. Calling it “the range” is shorthand, and the number 10^{-12} is a knob, not a constant of nature: it belongs in the error budget alongside V_CORE , which is the same kind of choice at the other end. Sec. 17.2 shows one consequence — because these thresholds are absolute while V rescales as $1/L^2$, the truncated problem is not exactly scale invariant even though the physics is.

With that understood, Eq. (9) is solved *on paper* for each shape:

$$\text{Pöschl-Teller: } \frac{v\mu^2}{\cosh^2(\mu R)} = \varepsilon \quad \implies \quad R = \frac{1}{\mu} \operatorname{arccosh} \sqrt{\frac{v\mu^2}{\varepsilon}} \quad (10)$$

$$\text{Gaussian: } v\mu^2 e^{-\mu^2 R^2} = \varepsilon \quad \implies \quad R = \frac{1}{\mu} \sqrt{\ln \frac{v\mu^2}{\varepsilon}} \quad (11)$$

$$\text{Lennard-Jones (tail): } \frac{1}{2} C_6 / R^6 = \varepsilon \quad \implies \quad R = \left(\frac{C_6}{2\varepsilon} \right)^{1/6} \quad (12)$$

$$\text{Lennard-Jones (core): } \frac{1}{2} C_{12} / r_{\min}^{12} = V_{\text{core}} \quad \implies \quad r_{\min} = \left(\frac{C_{12}}{2V_{\text{core}}} \right)^{1/12} \quad (13)$$

Why: An earlier version searched for these radii numerically, sweeping a global 40 000-point grid. Solving on paper removed that grid, removed a helper function, and made the result exact instead of limited by the sweep resolution.

4.2 The units of the coefficients

Since $[V_{\text{code}}] = \text{fm}^{-2}$ by Eq. (3), each term of the Lennard-Jones has to carry fm^{-2} on its own:

$$\left[\frac{C_{12}}{r^{12}} \right] = \text{fm}^{-2} \implies [C_{12}] = \text{fm}^{10}, \quad \left[\frac{C_6}{r^6} \right] = \text{fm}^{-2} \implies [C_6] = \text{fm}^4.$$

The same two exponents reappear in Sec. 14.1 as the powers of L in the scaling rule, and that is not a coincidence: a scaling law read off the units is the same statement twice.

For the other three potentials the strength v is dimensionless and the scale μ is an inverse length, $[\mu] = \text{fm}^{-1}$, so $v\mu^2$ carries the fm^{-2} by itself.

4.3 A factor of two in Eq. (121)

Equation (121) of [3] reads

$$V_{LJ}(r) = \frac{\hbar^2}{m_r} \left(\frac{C_{12}}{r^{12}} - \frac{C_6}{r^6} \right),$$

which in the units used here is $C_{12}/r^{12} - C_6/r^6$, with **no** $\frac{1}{2}$. Taken literally, it does not reproduce Table 4 of the same paper. Using the published deuteron parameters $C_{12} = 0.90485319$ and $C_6 = 6.81472$:

	a (fm)	r_0 (fm)
Table 4 of [3]	5.4	1.70
with the $\frac{1}{2}$	5.405	1.699
without it	1.435	1.412

Equations (70) and (120), for the well and the Gaussian, carry the same $\hbar^2/m_r$ prefactor and *do* reproduce Table 3 exactly as printed. The mismatch is therefore specific to Eq. (121); the likely reading is $2m_r$ in its denominator. The code uses the $\frac{1}{2}$.

Why: This is the kind of detail that survives unnoticed because the result *looks* right either way — the curves have the expected shape. Only the numerical comparison against the published table exposes it.

5 Numerov: the integrator

Numerov solves equations of the form $u'' = W(r)u$, which is exactly Eq. (4) with $W = 2(V - E)$. The recursion is

$$u_{i+1} = \frac{(12 - 10f_i)u_i - f_{i-1}u_{i-1}}{f_{i+1}}, \quad f_i = 1 - \frac{h^2 W_i}{12}. \quad (14)$$

Where it comes from

Expanding $u_{i\pm 1}$ in Taylor series and adding,

$$u_{i+1} + u_{i-1} = 2u_i + h^2 u_i'' + \frac{h^4}{12} u_i^{(4)} + O(h^6).$$

A second-order method simply discards the h^4 term. Numerov does not: since $u'' = Wu$, we have $u^{(4)} = (Wu)''$, and that second derivative can be approximated by the *same* central difference, without evaluating anything at new points. Substituting and regrouping gives Eq. (14).

Why: This is why Numerov reaches order h^4 at the cost of a second-order scheme: it never evaluates W away from the grid. The price is a *three-term* recursion — each point depends on the two before it, so an error made at the start travels all the way out.

```
f = 1.0 - h * h * W / 12.0
for i in range(1, POINTS - 1):
    w[i + 1] = ((12.0 - 10.0 * f[i]) * w[i] - f[i - 1] * w[i - 1]) / f[i + 1]
```

6 The logarithmic grid

6.1 The problem

Numerov requires $h\sqrt{|W|} \ll 1$ to represent the solution. Measured across the four potentials on a uniform grid:

potential	h (fm)	$\max W $	$h\sqrt{ W }$
square well	0.00103	8.3×10^{-1}	0.001
Pöschl-Teller	0.00851	2.1×10^0	0.012
Gaussian	0.00401	1.6×10^0	0.005
Lennard-Jones	0.00626	2.0×10^5	2.78

The Lennard-Jones is squeezed from both sides: its core reaches $V = 10^5$, so $\sqrt{2V} = 447$ and the characteristic length there is 0.002 fm, while its $1/r^6$ tail pushes R out to 125 fm. **No uniform step serves both.**

The consequence was measured, and it is a real error: the Lennard-Jones r_0 *drifted* as the number of points doubled, instead of converging.




Figure 2: Convergence test. On a uniform grid (red) r_0 never settles. On a logarithmic grid (blue) it locks in at the fifth decimal. A result that moves when the grid is refined is not a result.




Figure 3: Point distribution, same number of points in both cases. The uniform grid wastes almost everything on the tail and leaves the core with a handful of points; the logarithmic grid does the opposite.

6.2 The calculation

Strictly the argument of a logarithm has to be dimensionless, and r is in fm. So the substitution is

$$x = \ln\left(\frac{r}{r_*}\right), \quad r_* = 1 \text{ fm}, \quad (15)$$

and $r = r_* e^x$. Choosing $r_* = 1$ fm means every number below is unchanged — which is why the distinction is easy to drop — but the definition is only well posed with it written. Below we keep $r_* = 1$ implicit and write $r = e^x$.

With that, and $u = e^{x/2}w(x)$, using $\frac{d}{dr} = e^{-x} \frac{d}{dx}$:

$$u' = e^{-x} \frac{d}{dx} (e^{x/2}w) = e^{-x/2} \left(\frac{1}{2}w + w' \right), \quad (16)$$

$$u'' = e^{-x} \frac{d}{dx} \left[e^{-x/2} \left(\frac{1}{2}w + w' \right) \right] = e^{-3x/2} \left(w'' - \frac{1}{4}w \right). \quad (17)$$

Substituting into Eq. (4), $e^{-3x/2}(w'' - \frac{1}{4}w) = 2(V - E)e^{x/2}w$, so

$$w''(x) = \left[2(V(e^x) - E)e^{2x} + \frac{1}{4} \right] w(x) \quad (18)$$

Two things matter here. First, a $\frac{1}{4}$ **appears** that was not in the original equation — it comes from the substitution $u = e^{x/2}w$, not from the physics. Second, and decisively: Eq. (18) is still of the form $w'' = Ww$, which is the *only* thing Numerov requires. The recursion does not change at all; only what goes into it does.

Why: Changing the independent variable is the standard move when a problem has two very different scales. Here they differ by five orders of magnitude: 0.002 fm in the core against 125 fm in the tail.

```
r_start = pot.r_min if pot.r_min > 0.0 else pot.R * 1e-6
x = np.linspace(math.log(r_start), math.log(pot.R), POINTS)
h = x[1] - x[0]
r = np.exp(x)
W = 2.0 * (pot.V(r) - E) * r * r + 0.25
```

7 The two starts

Numerov needs *two* initial values, and the code picks between two cases:

```
if pot.r_min > 0.0:
    w[1] = h * math.exp(x[0] / 2.0)          # hard core: u = 0, u' = 1 there
else:
    w[0] = math.sqrt(r[0])                   # regular origin: u ~ r
    w[1] = math.sqrt(r[1])
```

Regular origin (well, Pöschl-Teller, Gaussian). Near $r = 0$ with finite V , the regular solution behaves as $u \simeq r$. Since $w = u e^{-x/2} = u/\sqrt{r}$, that gives $w \simeq \sqrt{r}$, which is exactly what the code writes.

Hard core (Lennard-Jones). There V blows up and u is exponentially small; we impose $u(r_{\min}) = 0$ with unit slope, and the same conversion $w = u/\sqrt{r}$ gives $w_1 \simeq h e^{x_0/2}$.

Why: Getting the start wrong is not a local mistake: by the three-term recursion, the error propagates through the whole integration. A measurement from this work: for the Lennard-Jones, the fourth-order start correction $u_1 = h + h^3 W_0/6$ amounts to 129% of h . It did not fix that version's problem — the cause was the grid — but it shows the scale of what is at stake.

8 The slope at the edge

Gluing the solution onto the analytic form needs $u(R)$ and $u'(R)$. The value comes out directly; the derivative uses a five-point one-sided difference:

$$\left. \frac{dw}{dx} \right|_R \simeq \frac{25w_N - 48w_{N-1} + 36w_{N-2} - 16w_{N-3} + 3w_{N-4}}{12h} + O(h^4), \quad (19)$$

and the conversion back to physical radius comes from the first line of the calculation in Sec. 6:

$$\left. \frac{du}{dr} \right|_R = e^{-x_N/2} \left(\frac{1}{2} w_N + \left. \frac{dw}{dx} \right|_R \right). \quad (20)$$

Why: Two reasons for the one-sided formula instead of a central difference. It is order h^4 , matching the rest of the method — a central one would be $O(h^2)$ and would spoil the global order. And it uses *interior points only*, so it never crosses the square well’s discontinuity.

The discontinuity finding

If a Numerov step straddles the well’s jump, the method assumes W is smooth across the three points it uses — and it is not. The convergence order was measured:

n	error in a	measured order
2001	8.9×10^{-4}	—
4001	4.5×10^{-4}	1.00
8001	2.2×10^{-4}	1.00
16001	1.1×10^{-4}	1.00

Order exactly 1, not 4. With the grid ending on the edge and the one-sided derivative, the same case drops to 1.9×10^{-10} — **six orders of magnitude**, with the same number of points.

9 What the code measures: a and r_0

9.1 The scattering length

Beyond R the solution is the straight line $u = C(1 - r/a)$. It crosses zero at $r = a$, and knowing $u(R)$ and $u'(R)$ fixes the line:

$$a = R - \frac{u(R)}{u'(R)} \quad (21)$$

Why: This is an *exact reading*, not a fit. No least squares, no extrapolation. The only source of error is the integration out to R .

The code carries a branch for $u'(R) = 0$: the line is horizontal, never crosses zero, and $a = \infty$.

```
if slope == 0.0:
    a = math.inf
else:
    a = pot.R - u[-1] / slope
```

Why: That branch is mathematically correct and **numerically unreachable**. A floating-point slope is never exactly zero. Swept across the square well’s resonance, it measures

v	a	$u'(R)$
1.2214	−80.3	7.9×10^{-3}
1.2325	−809	7.9×10^{-4}
1.2337	-1.82×10^6	3.5×10^{-7}
1.2349	+812	-7.8×10^{-4}
1.2460	+81.9	-7.8×10^{-3}

Small, then small and negative — never zero. Two things follow. First, $a = \infty$ never actually occurs, so unitarity has to be recognised by a threshold rather than by an equality. Second, and this is the part the branch hides: a **changes sign** through the resonance. Returning $+\infty$ would be choosing one side of that flip by fiat. The branch is kept only so the function is total.

What “infinite” means, numerically

The question is not rhetorical: at what value do we call a infinite? Asked about a itself the question has no answer, because a carries a length and the answer would then depend on the system. Asked about a dimensionless ratio it does. The right ratio is r_0/a : it is invariant under the rescaling of Sec. 14.1, and it vanishes at unitarity by construction. It is also the quantity Sec. 17.3 argues should be reported anyway.

R/a will not serve. R is a truncation we chose, not physics — and for the Lennard-Jones at unitarity $R/a = 1.5 \times 10^{-3}$, which is not small.

Measured over the twelve published cases:

	$ r_0/a $
at unitarity	2.2×10^{-11} to 2.1×10^{-5}
where an exterior node must count	3.14×10^{-1}

a gap of about 1.5×10^4 , so

$$\text{UNITARITY} = 10^{-3} \quad (22)$$

sits 47× above one group and 314× below the other. Every one of the twelve cases gets the Table 2 count for any threshold from 10^{-2} to 10^{-4} ; it fails at 10^{-5} , where the Lennard-Jones at unitarity stops being recognised. Three decades of margin, tested in `test_the_unitarity_threshold_is_not_fragile`. A threshold that worked at only one value would be a fit, not a criterion.

9.2 The normalisation, and why it is needed

Equation (4) is homogeneous: if u is a solution, so is Cu . The start chose an arbitrary normalisation, so u is not yet comparable with anything. We fix it by requiring that it meet the line at the edge:

$$u \longleftarrow u \frac{1 - R/a}{u(R)} \quad \Longrightarrow \quad u(R) = g_0(R), \quad g_0(r) \equiv 1 - \frac{r}{a}. \quad (23)$$

Why: Without this the effective-range integral would not be defined — it compares u against g_0 , and the comparison only makes sense once both are on the same scale.

9.3 The effective range

$$r_0 = 2 \int_0^\infty \left[g_0(r)^2 - u(r)^2 \right] dr. \quad (24)$$

The integrand measures *how far the true solution departs from the line*. Outside the range $u = g_0$ and it vanishes identically, so the integral does end at R .

Why: It is tempting to read r_0 as “the size of the potential”. For the square well the temptation is strong, because at every pole of a we get $r_0/R = 1$ exactly — and it is precisely because the analogy works so well there that it should not be generalised. In general r_0 is not a geometric radius:

- it can be **negative**;
- it depends on the shape of the *interior* wavefunction, not on where the potential stops;
- it can grow large near special structures, without the potential growing at all;
- it need not coincide with any radius one could draw on a plot of V .

What Eq. (24) says is narrower and safer: r_0 measures the failure of the true solution to be the straight line, weighted over all r . That is a statement about the wavefunction, and only indirectly about the potential.

Two details in the code:

The factor r . On the logarithmic grid $dr = r dx$, so the integral becomes $\int(\dots)r dx$. That is where the `* r` in the Simpson call comes from.

The analytic piece. The grid starts at r_{start} , not at zero. The stretch $[0, r_{\text{start}}]$ is added by hand: there u is either zero (inside a hard core) or proportional to r , and in both cases its contribution is $O(r_{\text{start}}^3)$ while the line's is $O(r_{\text{start}})$, leaving

$$2 \int_0^s g_0^2 dr = 2 \left(s - \frac{s^2}{a} + \frac{s^3}{3a^2} \right), \quad s \equiv r_{\text{start}}. \quad (25)$$

Why: This term used to apply only to the Lennard-Jones. Measurement showed the square well carried a residual error of 2.4×10^{-6} in r_0 that was exactly this missing piece. Making the term unconditional removed an `if` and improved precision by three orders of magnitude — down to 5.6×10^{-12} .

```
r0 = 2.0 * simpson((line ** 2 - u ** 2) * r, x=x)
s = r_start
r0 = r0 + 2.0 * (s - s ** 2 / a + s ** 3 / (3.0 * a ** 2))
```

10 Counting bound states

The code counts *zeros of the zero-energy solution*. It is a striking fact that the scattering solution, at $E = 0$, knows how many bound states the potential has.

Reference [3] treats this as a mandatory check rather than as a result: “*checking the number of nodes of the radial function is necessary to guarantee that it is indeed the situation we wanted to reproduce*”, expecting a nodeless $u(r)$ for $a < 0$ and one node for $a > 0$.

There is a subtlety. Beyond the edge the solution is the line $g_0 = 1 - r/a$, which crosses zero at $r = a$. If $a > R$ that zero is physical and counts. If a is effectively infinite, the zero sits at infinity and does not — hence the test

```
if pot.R < a and abs(r0 / a) > UNITARITY:
```

with `UNITARITY` fixed in Sec. 9.1. An earlier version read `R < a < 1e4`, a bound with a length hidden in it and no argument behind the number.

```
nodes = 0
for i in range(1, POINTS - 1):
    if u[i] * u[i + 1] < 0.0:
        nodes = nodes + 1
if pot.R < a < 1e4:
    nodes = nodes + 1
```

Why: The count is not optional. Two solutions can share (a, r_0) and differ in node count — and then they are *not* the same physical state. Every tuning ends by checking it.

A claim that used to live here, and is wrong: that the loop starts at `i = 1` because $u_0 = 0$ would otherwise count as a sign change. The test is `u[i]*u[i+1] < 0.0`, and a product of zero is not less than zero, so it never could. Running from `i = 0` and from `i = 1` gives the same count on all four potentials, and in three of them u_0 is not zero anyway. See the autopsy, Sec. 16, for what is and is not known about the original symptom.

11 The bound state

Beyond the edge, with $E < 0$, the solution must be the decaying exponential. That gives the matching condition

$$g(E) \equiv \frac{u'(R) + \kappa u(R)}{\max(|u(R)|, |u'(R)|, 1)} = 0, \quad \kappa = \sqrt{-2E} \quad (26)$$

Why: Two separate decisions are packed into that expression.

The sum, not the ratio $u'/u + \kappa$. Measured: with the ratio, `brentq` fails to bracket for three of the four potentials — it reports that the two ends of the interval carry the same sign. The sum brackets for all four. (An earlier version of these notes justified this by saying that $u(R)$ passes through zero because “the deuteron has a node”. That is wrong twice over. The deuteron’s bound state is the ground state and has *no* node — what carries a node is the zero-energy scattering solution of a potential that supports a bound state, which is a different function. And $|u(R)|$ never comes near zero here: swept across the bracket, its smallest value over the four potentials is 7×10^{-3} .)

The denominator. It is normalisation, and it does not move the root — removing it changes the answer by 3×10^{-16} . What it fixes is scale. The raw residual $u'(R) + \kappa u(R)$ is of order 10^{-2} for the square well and 10^{21} for the Lennard-Jones, because u at the edge spans that range after climbing out of the hard core. A single absolute tolerance cannot serve both. Dividing brings every potential to order unity.

Why there is no search

The finite-range formula already delivers the answer to about 1%. It is the guess, and `brentq` merely refines it within $[1.8g, 0.3g]$:

```
return brentq(gap, 1.8 * guess, 0.3 * guess, xtol=1e-15)
```

Why: An earlier version scanned 120 energies on a logarithmic grid just to find where the function changed sign. That was waste: the approximate answer was already in hand and was being ignored. Replacing the scan by the guess made the calculation $3.5\times$ faster *and* more accurate. There is also a conceptual bonus: the bracket always working is itself a statement that the effective-range expansion is good.

12 The inverse problem: tuning

Given a target (a, r_0) , find the parameters. It is a nonlinear 2×2 system, and a has *poles*: it jumps from $+\infty$ to $-\infty$ every time a new bound state appears.

12.1 Two nested loops

<code>tune</code>	outer loop: moves the scale until r_0 matches
<code>match_a</code>	inner loop: moves the strength until $1/a$ matches
<code>find_root</code>	widens a bracket until the sign flips, then bisection

Each loop solves *one* equation in *one* variable. Inside the outer loop, `r0_error` runs the entire inner loop again for every scale it tries — which is why this is the slow part, about 35 of the 54 seconds.

The decision that makes it work: chase $1/a$

Why: a has poles; $1/a$ is smooth and crosses zero. Bisection *dies* on a pole and thrives on a zero. As a bonus, unitarity becomes $1/a = 0$ and needs no special case at all.

```
if math.isinf(a_target):
    inv_a_target = 0.0
else:
    inv_a_target = 1.0 / a_target
```

Why not a 2D Newton

The level sets of $1/a$ and r_0 in the (strength, scale) plane are *nearly parallel* over much of the domain. A Newton solver on both equations at once would meet an ill-conditioned Jacobian exactly there, and take large steps in the wrong direction. Nested bisection cannot diverge once the sign is bracketed — slower per iteration and far more reliable for code someone else will run.

The six return values

`tune` returns (strength, scale, a, r0, nodes, ok). The last is a flag: false if r_0 failed to converge *or* if the node count missed its target. Without it, a tuning that landed on the wrong state would pass unnoticed.

Why: What comes out of `tune` is a different pair of numbers for every potential — there is no shared parameter. What is shared is the pair (a, r_0) , and with it everything a low-energy measurement can see. That is the point of Sec. 1.

13 Every constant, one by one

No number in the code is a guess. This table is the justification for each.

constant	value	where	justification
V_ZERO	10^{-12}	defines R	Below this the potential is treated as gone. Measured: cutting at 10^{-8} truncates the Lennard-Jones tail and misses r_0 by 0.16%; 10^{-12} closes to 0.004%.
V_CORE	10^5	LJ start	Where integration begins inside the core. Ref. [3] prescribes $V(r_{\min}) \sim 10^{10}$; we measured that 10^3 breaks the physics (r_0 off by 0.2%) while anything from 10^4 up changes nothing. Starting deeper only costs stiffness.
POINTS	8001	grid size	Total error is truncation (falls with n) plus roundoff (grows with n); 8001 sits at the minimum. The three purely attractive potentials are at 10^{-11} and do not care. The Lennard-Jones sets the value: past ~ 8000 its r_0 <i>degrades</i> again — 3×10^{-5} at 32001, 3×10^{-4} at 64001.
r_start	$R \times 10^{-6}$	regular origin	Where the logarithmic grid begins when there is no core. It cannot be 0 because $\ln 0$ does not exist. The stretch below it enters analytically.
[1.8g, 0.3g]	—	brentq	Bracket around the finite-range guess. It works on both systems tested; that it works at all is a statement about the quality of the expansion.
1.3	—	find_root	Bracket widening factor, up to 40 attempts. If no sign change is found it raises rather than returning garbage.
$10^{-7}, 10^{-6}$	—	tolerances	Convergence of r_0 in tune and in the tests. Measured: the tuning hits its target to 10^{-10} , comfortably inside them.

14 The tabulated data

Five dictionaries, all traceable:

TARGETS	Table 2 of [3]: the three targets. The bound-state count is <i>not</i> a column of that table; it comes from the text of Sec. 4.5.
PUBLISHED	Tables 3 and 4: the published parameters. Used as the initial guess <i>and</i> as the validation target.
SYSTEMS	Table 1: two real systems, with measured energies.
GUESS	The deuteron solution, used as a starting point for other systems.
SOURCES	The three references with DOI, plus the D1 divergence.

14.1 Scale invariance in GUESS

The problem is scale invariant: multiplying all lengths by L takes $(a, r_0) \rightarrow (La, Lr_0)$. So attacking another system only needs the guess rescaled — but the exponents follow from the units, not from the powers of r alone. Since $[V_{\text{code}}] = \text{fm}^{-2}$, the rescaled potential must satisfy $V'(Lr) = V(r)/L^2$. For the repulsive term,

$$\frac{C'_{12}}{(Lr)^{12}} = \frac{1}{L^2} \frac{C_{12}}{r^{12}} \implies C'_{12} = L^{10} C_{12},$$

and the same argument on the attractive term gives $C'_6 = L^4 C_6$. So

$$\mu \rightarrow \frac{\mu}{L}, \quad C_6 \rightarrow C_6 L^4, \quad C_{12} \rightarrow C_{12} L^{10}. \quad (27)$$

These are exactly $[C_6] = \text{fm}^4$ and $[C_{12}] = \text{fm}^{10}$ read backwards. An earlier version of these notes carried $C_{12} \rightarrow C_{12} L^6$, which comes from treating V as dimensionless; it is wrong, and it misses C_6 entirely. The consequence was not cosmetic — see Sec. 17.2.

14.2 The constant that is tabulated nowhere

$\hbar^2/2m_r$ appears in no table. We obtain it by inverting the zero-range formula on the published pair (a, E_{zr}) :

$$E_{zr} = -\frac{\hbar^2/2m_r}{a^2} \implies \frac{\hbar^2}{2m_r} = -E_{zr} a^2. \quad (28)$$

system	inferred	literature	deviation
deuteron	41.462 MeV·fm ²	41.47	0.02%
⁴ He dimer	12.095 K·Å ²	12.12	0.21%

Why: This is a **test**, not a convenience. Recovering both known constants proves Table 1 is internally consistent — and, incidentally, that the factor of 2 in Eq. (5) is right.

The three tunings that do not match, and why

Nine of the twelve tunings land within 0.2% of the published parameters. Three do not, and they do not fail for the same reason.

Two of them are the D1 divergence. (nn, Lennard-Jones) at 2.68% and (deuteron, Pöschl-Teller) at 2.65%. The cause is *not* numerical: refining the grid moves the deviation from 2.70% to 2.69%. Table 3 of [3] reports $r_0 = 1.73$ for the mPT deuteron row and Table 4 reports $r_0 = 2.71$ for the Lennard-Jones nn row, while Table 2 of the same paper lists the targets as 1.70 and 2.70. Feeding the published parameters back through the solver confirms it: they produce $r_0 = 1.7300$ and 2.7080, not the targets. We tune to the target and land elsewhere — correctly.

The third is a flat direction, not a disagreement. (deuteron, Lennard-Jones) sits 1.13% away, but here the published parameters *do* reproduce the target: they give $r_0 = 1.6990$ against 1.7000, an error of 0.06%. Both parameter pairs are therefore valid. What happens is that the two Lennard-Jones coefficients are strongly correlated. Starting from the published pair and moving to ours:

change	effect on r_0
C_6 alone, +0.47%	−1.09%
C_{12} alone, +1.13%	+1.18%
both together	+ 0.05%

Each coefficient on its own moves r_0 by about 1%; together they cancel. The inversion $(a, r_0) \rightarrow (C_{12}, C_6)$ has a nearly flat direction, so the observable is pinned down far more tightly than the parameters that produce it. This is worth stating plainly because the naive reading — “the parameters disagree by 1%, so something is wrong” — is exactly backwards. Note also that Lennard-Jones is the *most* sensitive of the four potentials ($d \ln r_0 / d \ln C_6 = 2.28$, against 0.42 for the well), so the flatness comes from the correlation between the two coefficients, not from any insensitivity of the potential.

General rule: a discrepancy that survives refinement is never numerical.

15 The checks

`test_lab.py` runs eight checks, and every one compares against something the code did **not** compute:

check	external anchor
a of the well	Eq. (80) of [3], closed form
r_0 of the well	Eq. (92) of [3], closed form
$r_0/R = 1$ at the poles	exact prediction, caption of Fig. 6 of [3]
bound state	analytic condition $k \cot(kR) = -\kappa$
node counts	Table 2, all 12 cases
tuning	Tables 3 and 4, all 12 cases
inferred constants	41.47 and 12.12 from the literature
grid convergence	halving the points may change nothing

Why: No check says “same as last time”. Such a check would only prove the code is consistent with its own mistakes.

The last one is what *found* the uniform-grid bug. It halves the points rather than doubling them, deliberately: since the total error has a minimum, the useful question is “have we converged coming from below?” and not “what happens deep in the roundoff regime?”.

16 Autopsy: the errors made

This section exists because the cause of an error is usually worth more than the result.

1. Numerov dropping from order 4 to order 1

Symptom: the square well converged far too slowly. **Cause:** a step straddled the discontinuity; Numerov assumes W is smooth across the three points it uses. **Fix:** the grid ends exactly on the edge and the derivative uses interior points only. **Gain:** six orders of magnitude.

2. The Lennard-Jones r_0 drifting

Symptom: $1.74257 \rightarrow 1.74215 \rightarrow 1.74165$ as the points doubled. **Diagnosis:** two wrong hypotheses were discarded before the right one — it was not the tail (truncating at 30 fm kept the drift) and not the start (a fourth-order correction changed nothing). It was $h\sqrt{|W|} = 2.78$ in the core. **Fix:** logarithmic grid. **Confirmation:** the new value agrees with an independent method to seven digits.

3. Matching inside the well

Symptom: the well's energy came out 0.44% wrong. **Cause:** the match used the *last* point where the potential still acted — which, in a well, is inside, where $V \neq 0$. **Fix:** match outside.

4. The node count, and a diagnosis that was wrong

Symptom: every potential reported one extra bound state. **Recorded cause:** $u_0 = 0$ counted as a sign change. **That diagnosis does not survive checking.** The test in the code is $u[i] * u[i+1] < 0.0$, and a product of zero is not less than zero, so $u_0 = 0$ could never have been counted. Running the loop from $i = 0$ and from $i = 1$ gives the *same* count for all four potentials, and in three of them u_0 is not even zero. Starting at $i = 1$ is harmless, but it cannot have been what fixed the symptom. What did is no longer recorded, and is deliberately not reconstructed: an autopsy with a gap in it is more use than one with a guess in it.

5. The inner term that only applied to the Lennard-Jones

Symptom: residual error of 2.4×10^{-6} in the well's r_0 . **Cause:** the logarithmic grid does not cover $[0, r_{\text{start}}]$, and the analytic term sat inside an **if**. **Fix:** make it unconditional — one **if** fewer and three orders of magnitude better.

6. Citations taken on trust

Symptom: three potentials were attributed to the wrong equation numbers (74 instead of 70, 118 instead of 120, 120 instead of 121), and the node counting was attributed to a theorem that the paper never mentions. **Cause:** the citations were written from memory rather than checked against the source. **Fix:** a full pass with the paper open, which also turned up the factor of two in Sec. 4.3. **Lesson:** a plausible citation and a verified citation look identical on the page.

17 Results




Figure 4: Three potentials tuned to the same (a, r_0) of the deuteron: the wavefunctions are *identical outside* the range and *completely different inside*. Dotted lines mark where each potential ends. Everything a low-energy experiment can see lives outside.




Figure 5: Reproduction of Figs. 5 and 6 of [3]. Left: a/R diverges at $\sqrt{2}v_0 = \pi/2 + n\pi$, where a new bound state appears. Right: at each of those points $r_0/R = 1$ *exactly* (red stars), as stated in the caption of Fig. 6 of [3]. Orange lines mark the *other* family of singularities, the roots of $\tan x = x$, where $a = 0$ and it is r_0 that diverges.

17.1 Binding energies

method	deuteron ($r_0/ a = 0.32$)		^4He dimer ($r_0/ a = 0.089$)	
	E (MeV)	deviation	E (K)	deviation
experiment	-2.224	—	-1.620×10^{-3}	—
zero range	-1.416	-36.3%	-1.480×10^{-3}	-8.6%
finite range	-2.223	-0.05%	-1.628×10^{-3}	+0.47%
square well	-2.204	-0.91%	-1.627×10^{-3}	+0.46%
Pöschl-Teller	-2.227	+0.11%	-1.628×10^{-3}	+0.47%
Gaussian	-2.214	-0.47%	-1.627×10^{-3}	+0.46%
Lennard-Jones	-2.203	-0.92%	—	—

Two systems separated by 13 orders of magnitude in energy, described by the same two numbers and the same two formulas. That is the universality of Sec. 1, measured.

The residual spread of about 1% *among* the potentials is the shape dependence — the $O(k^4)$ terms, the only thing still distinguishing a square well from a Lennard-Jones once a and r_0 are fixed.

17.2 The helium Lennard-Jones, and what the wrong exponent cost

The scaling rule of Sec. 14.1 is not decoration: it supplies the starting guess that the tuning begins from. With the wrong exponent the guess for the helium dimer was off by a factor of 443 in *both* coefficients, and — worse — it landed on the branch with **zero** bound states when the target has one:

starting guess	C_6	C_{12}	what it gives
$C_{12}L^6, C_6$ untouched	7.84	1.19×10^4	$a = 1.76, r_0 = 1.33, 0$ nodes
$C_6L^4, C_{12}L^{10}$	3476	5.27×10^6	$a = 24.8, r_0 = 7.995, 1$ node
target	$a = 90.4, r_0 = 8.0, 1$ node		

The correct rule arrives with r_0 already within 0.06% of the target and on the right branch. The wrong one does not converge at all: it was left running for 170 s without finishing, while the corrected guess converges in 83 s.

That non-convergence is the reason the Lennard-Jones was quietly absent from the helium row — three potentials were run there, not four, and no note in these notes said so. With the exponents fixed the case runs, and it belongs with the others:

$$E(^4\text{He}_2, \text{Lennard-Jones}) = -1.627 \text{ mK}, \quad \text{measured} - 1.620 \text{ mK}.$$

The lesson is about the check, not the exponent. Scale invariance is a *property* the code must satisfy, so it can be tested directly: rescale every length by L , and (a, r_0) must come back as (La, Lr_0) . That test is three lines, needs no reference value, and would have caught this immediately. It is now in `test_lab.py`.

17.3 Accuracy achieved

check	error	kind
a of the well against Eq. (80)	1.9×10^{-10}	relative
r_0 of the well against Eq. (92)	5.6×10^{-12}	relative
bound state against $k \cot(kR) = -\kappa$	3.4×10^{-10}	relative
r_0/R at the four poles	1.0000000000	exact value, not an error
solver against closed forms (16 points)	8.4×10^{-9}	relative
tuning, a against target	10^{-6}	relative
tuning, r_0 against target	10^{-6}	absolute

Why: The last two rows used to be one row reading “ 10^{-10} ”. Splitting them is not pedantry: the code asserts `abs(r0 - target) < 1e-6`, an absolute tolerance in fm, and `abs(a/target - 1) < 1e-6`, a relative one. Those are different requirements, and at unitarity, where $a \rightarrow \infty$, only the absolute one is even meaningful. Stating “the error is 10^{-6} ” without saying which is an invitation to compare two numbers that are not comparable.

Convention for the rest of this document: every error is relative unless it is marked absolute.

The dimensionless number worth reporting

Absolute and relative are the minimum. Better, where the physics offers one, is a ratio built from the problem itself. Here the natural one is r_0/a : it is dimensionless, it is invariant under the rescaling of Sec. 14.1, and it measures how close a system sits to unitarity, where it vanishes.

system	a	r_0	r_0/a	zero-range error
unitarity	∞	1.0	0	—
^4He dimer	90.4	8.0	0.089	8.6%
neutron-neutron	-18.5	2.7	-0.146	—
deuteron	5.4112	1.744	0.322	36.3%

Why: This single column explains a result that otherwise looks arbitrary. The zero-range formula misses the deuteron by 36% and the helium dimer by only 8.6% — a factor of 4.2. The ratio r_0/a differs between them by a factor of 3.6. Zero range is the approximation that throws r_0 away, so its error is governed by exactly this number, and the two factors agreeing to within 20% is the expansion behaving as advertised. Note also that r_0/a is negative for the neutron-neutron system, which is a sign that no dimensionless “size” reading of r_0 survives: see Sec. 9.

18 The error budget

Every number in this document rests on six choices that are ours, not the physics’. This section varies each one and reports what r_0 does. All entries are **relative** errors against the value at the setting actually used, and the deuteron parameters of Table 3 are the test case throughout.

POINTS — discretisation and round-off

POINTS	well	Pöschl-Teller	Gaussian	Lennard-Jones
2001	2.1×10^{-9}	1.6×10^{-10}	8.3×10^{-11}	2.6×10^{-7}
4001	1.3×10^{-10}	2.0×10^{-12}	3.9×10^{-12}	3.7×10^{-7}
8001	—	—	—	—
16001	8.2×10^{-12}	5.0×10^{-11}	1.9×10^{-11}	1.9×10^{-6}
32001	1.4×10^{-11}	2.7×10^{-10}	1.3×10^{-10}	3.4×10^{-5}

Why: This is the one knob where more is not better, and the table shows why. The three purely attractive potentials are flat from 4001 upward — they sit at machine precision and stop caring. The Lennard-Jones does not: its error falls to 8001 and then *climbs* by two orders of magnitude out to 32001. That is round-off overtaking truncation, and it is the reason 8001 is written down rather than “as many as we can afford”. The review asked for values beyond 4001 and 8001; the answer is that they are worse.

V_ZERO — truncating the tail

V_ZERO	well	Pöschl-Teller	Gaussian	Lennard-Jones
10^{-9}	0	1.4×10^{-8}	1.6×10^{-9}	1.1×10^{-3}
10^{-10}	0	1.5×10^{-9}	2.1×10^{-10}	6.4×10^{-4}
10^{-12}	—	—	—	—
10^{-14}	0	3.2×10^{-10}	4.7×10^{-11}	3.6×10^{-4}

Why: The well is exactly zero in every row, which is the point made under Eq. (9): it is the only potential with a real range, so truncation cannot touch it. Everything else pays, and the Lennard-Jones pays 10^3 to 10^6 times more than the others, because a $1/r^6$ tail is still there when a Gaussian has long since vanished. Note that 10^{-14} is also worse than 10^{-12} : pushing the cut further out lengthens the grid and buys round-off.

V_CORE — a wall instead of the core

V_CORE	well	Pöschl-Teller	Gaussian	Lennard-Jones
10^3	0	0	0	2.4×10^{-3}
10^4	0	0	0	2.9×10^{-7}
10^5	—	—	—	—
10^6	0	0	0	6.7×10^{-7}
10^7	0	0	0	3.2×10^{-7}

Why: Only the Lennard-Jones has a core, so only it responds — the zeros in the other three columns are exact, not rounded. From 10^4 upward the answer is converged to 10^{-6} , and 10^3 is visibly wrong. This is the measurement behind using 10^5 where [3] prescribes 10^{10} : the extra five decades buy nothing and cost stiffness.

R_START_FRACTION — skipping the origin

fraction of R	well	Pöschl-Teller	Gaussian	Lennard-Jones
10^{-3}	2.9×10^{-9}	$\mathbf{6.8 \times 10^{-6}}$	4.5×10^{-7}	0
10^{-4}	1.8×10^{-12}	6.7×10^{-9}	4.5×10^{-10}	0
$\mathbf{10^{-6}}$	—	—	—	—
10^{-8}	8.5×10^{-12}	1.5×10^{-11}	9.6×10^{-12}	0
10^{-10}	5.7×10^{-11}	1.1×10^{-11}	4.4×10^{-12}	0

Why: The Lennard-Jones column is exactly zero everywhere because it never uses this knob: it has a hard core, so its start comes from `V_CORE`. For the other three the piece below r_{start} is added analytically, which is why even 10^{-3} costs only 10^{-6} . Converged by 10^{-4} ; the value used has four decades of margin.

This constant used to be written inline inside `numerov`. It was lifted out to module level while assembling this table — a knob that cannot be turned cannot be audited.

The edge derivative — h^4 against h^2

Replacing the five-point one-sided stencil by the three-point one, and reading the change in a :

potential	change in a
square well	3.7×10^{-6}
Pöschl-Teller	3.5×10^{-7}
Gaussian	9.2×10^{-8}
Lennard-Jones	5.0×10^{-7}

Why: The square well is hit ten times harder than the others, and that is not an accident: it is the one with a discontinuity at the edge, so the stencil is sampling a function whose derivative jumps. Dropping to $O(h^2)$ would put the matching error at 10^{-6} , four orders above the 1.9×10^{-10} that Sec. 15 reports against Eq. (80). The whole accuracy of this laboratory rests on those two extra stencil points.

brentq — locating the root

Deuteron binding energy, square well, as the tolerance is loosened:

<code>xtol</code>	E	change
10^{-6}	-0.02628469	4.0×10^{-6}
10^{-9}	-0.02628459	2.3×10^{-10}
10^{-12}	-0.026284587	3.7×10^{-12}
$\mathbf{10^{-15}}$	-0.026284587	—

Why: Cheap: the root finder is not where the error lives. Even 10^{-9} would leave the answer good to 10^{-10} , which is below every other term in this budget. `xtol` is set to 10^{-15} because there is no reason not to.

What the budget says

source	dominant for	size at the settings used
V_ZERO	Lennard-Jones	$\lesssim 10^{-4}$ if loosened
POINTS	Lennard-Jones	10^{-5} if pushed to 32001
edge derivative	square well	10^{-6} if dropped to h^2
V_CORE	Lennard-Jones	10^{-7}
R_START_FRACTION	Pöschl-Teller	10^{-11}
brentq	—	10^{-12}

Three things are worth taking away. The Lennard-Jones dominates four of the six rows, because a hard core and a power-law tail stress every truncation at once. The square well is immune to the truncation knobs and uniquely sensitive to the edge derivative, for the same reason in both cases: it is the only potential with a genuine boundary. And no single knob is close to limiting the others — which is what one wants from a budget, and is why the review’s judgement that a full study “is not the case here” is fair. What is written above is the map, not the survey.

19 A potential with no free parameters

Everything so far has been a shape with two knobs, turned until it sat at a chosen (a, r_0) . The Aziz HFD-B potential [2] is not that. Its parameters are fixed by spectroscopy and *ab initio* theory, so feeding it in and reading (a, r_0) out is a **prediction**. It is the only falsifiable object in this document.

$$V(r) = \varepsilon \left[A e^{-\alpha x + \beta x^2} - F(x) \left(\frac{C_6}{x^6} + \frac{C_8}{x^8} + \frac{C_{10}}{x^{10}} \right) \right], \quad x = \frac{r}{r_m},$$

with $F(x) = e^{-(D/x-1)^2}$ for $x < D$ and 1 otherwise.

Two things it forced

The units had to be used, not assumed. The parameters are published in kelvin and ångström. Getting them into the solver means $V_{\text{code}} = V_{\text{phys}}/(2\hbar^2/2m_r)$ from Eq. (6), with $\hbar^2/2m_r = 12.09 \text{ K}\cdot\text{\AA}^2$; lengths then come out in ångström and V_{code} in $\text{\AA}^{-2}$. This is the first place where the convention of Sec. 2 does work rather than sit there, and it checks out: the potential reproduces its own tabulated minimum, $V(r_m) = -\varepsilon = -10.948 \text{ K}$, to six digits.

The core is soft. The Lennard-Jones diverges as r^{-12} , so the code starts outside it at $V = V_{\text{CORE}}$. Aziz does not diverge. As $r \rightarrow 0$ the exponential saturates and $F(x)$ switches the dispersion terms off, leaving a ceiling

$$V_{\text{max}} = \frac{\varepsilon A}{2\hbar^2/2m_r} = 8.35 \times 10^4,$$

which is *below* $V_{\text{CORE}} = 10^5$. There is no radius where $V = V_{\text{CORE}}$, so $r_{\text{min}} = 0$ and the regular-origin start applies. Asking for the hard-core radius raises “**f(a) and f(b) must have different signs**” — the bracket reporting, correctly, that the thing being looked for does not exist.

What it predicts, and what that checks

	this work	published HFD-B [1]	deviation
a	88.434 Å	88.430 Å	0.005%
r_0	7.276 Å	7.276 Å	0.00%
E	−1.692 mK	−1.69 mK	0.1%
nodes	1	1	—

This is the strongest check in this document. The four tunable potentials are fitted to (a, r_0) , so reproducing (a, r_0) with them is circular. Aziz is not fitted to anything here, so agreeing to four digits with an independent calculation using the same potential tests the integrator, the matching, the effective-range integral, the bound-state search and the unit conversion simultaneously, against a number this work had no part in choosing.

Note also that Table 1 of [3] quotes 90.4 Å and −1.62 mK. Those are the Cencek *et al.* values [5], a 2012 *ab initio* surface, not HFD-B. The 2% gap between the two columns is a real difference between two potentials, and both are reproduced by the same solver.

The constant matters here, and only here

Getting that agreement required using $\hbar^2/2m_r = 12.1193 \text{ K}\cdot\text{Å}^2$ from CODATA rather than the 12.0948 that Sec. 14 *infers* by inverting the zero-range formula. The two differ by 0.20%. With the inferred value, a comes out 82.98 Å: 6% low.

That factor of thirty amplification is the physics, not a defect:

$$\frac{d \ln a}{d \ln(\hbar^2/2m_r)} = 48.7 \quad (29)$$

near unitarity. The helium dimer is a halo state — $|a|/r_0 = 12$ — which is another way of saying it sits close to a pole of a , and everything near a pole is amplified.

Why: Two consequences worth carrying forward. **First**, this is why the inferred constant is good enough everywhere else: the four tunable potentials are fitted to a target, so a small error in the constant is absorbed into the fitted parameters and never shows. Only a potential with no free parameters has nowhere to hide it. **Second**, Eq. (29) is a warning label for the three-body work: if 0.2% in a constant becomes 6% in a for two particles, quantities computed near unitarity need their inputs specified to more digits than feels reasonable.

20 What is missing

No error bars. No number here carries an estimated discretisation uncertainty.

The Aziz HFD-B potential is now implemented — see Sec. 19. What is still missing there is a published HFD-B result to check against; see that section for what could and could not be verified.

No variable phase method. An independent second route to a , which would give a cross-check on the least reliable number here.

No finite- k phase shift. Everything is done at $E = 0$, with r_0 from the integral. Computing $\delta_0(k)$ and fitting Eq. (1) would be an independent validation — and is the test that would settle the Lennard-Jones r_0 , reliable only to $\sim 10^{-6}$ and limited by roundoff.

s -wave only, valid while $kR \ll 1$.

References

- [1] P. Stipanović, L. Vranješ Markić and J. Boronat, *Elusive structure of helium trimers*, J. Phys. B **49**, 185101 (2016), arXiv:1607.05872. Table 1, row HFDB.
- [2] R. A. Aziz, F. R. W. McCourt and C. C. K. Wong, *A new determination of the ground state interatomic potential for He₂*, Molec. Phys. **61**, 1487 (1987). Parameters as tabulated in J. Boronat and J. Casulleras, arXiv:cond-mat/9309015, Appendix A.
- [3] M. Macêdo-Lima and L. Madeira, *Scattering length and effective range of microscopic two-body potentials*, Rev. Bras. Ensino Fís. **45**, e20230079 (2023). doi:10.1590/1806-9126-RBEF-2023-0079
- [4] R. W. Hackenburg, *Low-energy neutron-proton effective range parameters*, Phys. Rev. C **73**, 044002 (2006). doi:10.1103/PhysRevC.73.044002
- [5] W. Cencek *et al.*, *Effects of adiabatic, relativistic, and QED corrections on the pair potential of helium*, J. Chem. Phys. **136**, 224303 (2012). doi:10.1063/1.4712218

A The three functions in full

The body of the central functions, so that no line is left unexplained. The numbered comments refer to the notes below each block.

A.1 numerov — the integrator

```
def numerov(pot, E):
    r_start = pot.r_min if pot.r_min > 0.0 else pot.R * 1e-6           # (1)

    x = np.linspace(math.log(r_start), math.log(pot.R), POINTS)       # (2)
    h = x[1] - x[0]
    r = np.exp(x)
    W = 2.0 * (pot.V(r) - E) * r * r + 0.25                           # (3)
    f = 1.0 - h * h * W / 12.0                                         # (4)

    w = np.zeros(POINTS)
    if pot.r_min > 0.0:
        w[1] = h * math.exp(x[0] / 2.0)                               # (5)
    else:
        w[0] = math.sqrt(r[0])                                         # (6)
        w[1] = math.sqrt(r[1])
    for i in range(1, POINTS - 1):
        w[i + 1] = ((12.0 - 10.0*f[i]) * w[i] - f[i-1]*w[i-1]) / f[i+1] # (7)

    u = np.exp(x / 2.0) * w                                           # (8)

    dwdx = (25.0*w[-1] - 48.0*w[-2] + 36.0*w[-3]
            - 16.0*w[-4] + 3.0*w[-5]) / (12.0 * h)                    # (9)
    slope = math.exp(-x[-1] / 2.0) * (0.5 * w[-1] + dwdx)            # (10)

    return r, u, x, h, slope, r_start
```

1. Where to start: outside the core if there is one, otherwise near the origin, because $\ln 0$ does not exist. The skipped stretch enters analytically.
2. The grid is uniform in $x = \ln(r/r_*)$, $r_* = 1$ fm by Eq. (15), and ends *exactly* at $\ln R$.
3. The effective potential of Eq. (18). The r^2 comes from e^{2x} ; the 0.25 is the $\frac{1}{4}$ the substitution generated.

4. Numerov's f . Note it is a *vector*, computed once for the whole grid before the loop.
5. Hard-core start: $u = 0$, $u' = 1$, converted through $w = u/\sqrt{r}$.
6. Regular start: $u \simeq r$, hence $w \simeq \sqrt{r}$. Two points are needed because the recursion has three terms.
7. The recursion. It is the only line that actually integrates the equation.
8. Back from w to u . From here on everything is in physical units.
9. Five-point one-sided derivative, order h^4 , interior points only.
10. Conversion of the derivative to radius: $du/dr = e^{-x/2}(w/2 + dw/dx)$.

A.2 scattering — the three readings

```
def scattering(pot):
    r, u, x, h, slope, r_start = numerov(pot, 0.0) # (1)

    if slope == 0.0:
        a = math.inf # (2)
    else:
        a = pot.R - u[-1] / slope

    u = u * (1.0 - pot.R / a) / u[-1] # (3)
    line = 1.0 - r / a

    r0 = 2.0 * simpson((line ** 2 - u ** 2) * r, x=x) # (4)
    s = r_start
    r0 = r0 + 2.0 * (s - s**2 / a + s**3 / (3.0 * a**2)) # (5)

    nodes = 0
    for i in range(1, POINTS - 1): # (6)
        if u[i] * u[i + 1] < 0.0:
            nodes = nodes + 1
    if pot.R < a < 1e4: # (7)
        nodes = nodes + 1

    return a, r0, nodes
```

1. Zero energy: the only case in which the exterior solution is a straight line.
2. A horizontal line means exact unitarity, and $a = \infty$ is the right answer — not an error to be avoided.
3. The normalisation. Without it the following integral would not be defined, because the equation is homogeneous.
4. The integral of Eq. (24). The $* r$ is the Jacobian $dr = r dx$.
5. The stretch $[0, r_{\text{start}}]$, done on paper.
6. The loop starts at 1. This is cosmetic, not a fix: the test is $u[i]*u[i+1] < 0$, and a zero product never counts, so starting at 0 gives the same answer — verified on all four potentials.
7. The exterior node at $r = a$, which exists only if a is finite and larger than R .

A.3 tune — the inverse problem

```
def tune(name, a_target, r0_target, strength, scale, nodes_target=None):
    if math.isinf(a_target):
        inv_a_target = 0.0 # (1)
```

```

else:
    inv_a_target = 1.0 / a_target

for step in range(25):
    strength = match_a(name, scale, strength, inv_a_target)
    a, r0, nodes = scattering(POTENTIALS[name](strength, scale))
    if abs(r0 - r0_target) < 1e-7:
        break

    def r0_error(trial_scale):
        s = match_a(name, trial_scale, strength, inv_a_target)
        a2, r02, n2 = scattering(POTENTIALS[name](s, trial_scale))
        return r02 - r0_target

    scale = find_root(r0_error, scale)

    strength = match_a(name, scale, strength, inv_a_target)
    a, r0, nodes = scattering(POTENTIALS[name](strength, scale))

    ok = abs(r0 - r0_target) < 1e-6
    if nodes_target is not None and nodes != nodes_target:
        ok = False

    return strength, scale, a, r0, nodes, ok

```

1. Unitarity with no special case: $a = \infty$ is simply $1/a = 0$.
2. Twenty-five attempts at most. Leaving early means it converged.
3. Inner loop: fixes the strength so that $1/a$ matches, scale held.
4. This nested function is the cost of the algorithm: for every scale it tries, it runs the *entire* inner loop again.
5. One last adjustment of the strength, because the scale moved last.
6. The `ok` flag. Without it, a tuning that landed on the wrong state would pass unnoticed.

B The data, with values

TARGETS — Table 2 of [3]

case	a (fm)	r_0 (fm)	bound states
nn (neutron-neutron)	-18.5	2.7	0
unitarity	$\pm\infty$	1.0	0
deuteron	5.4	1.7	1

The bound-state column is not printed in Table 2; it comes from the text of Sec. 4.5 of [3], which asks for a nodeless $u(r)$ when $a < 0$ and one node when $a > 0$.

PUBLISHED — Tables 3 and 4, against what the tuning returns

Throughout, p_1 is the **strength** parameter and p_2 the **scale** parameter, in that order — which is the order the constructors take them in. Concretely:

potential	p_1 (strength)	p_2 (scale)
square well	v , dimensionless	μ , fm ⁻¹
Pöschl-Teller	v , dimensionless	μ , fm ⁻¹
Gaussian	v , dimensionless	μ , fm ⁻¹
Lennard-Jones	C_6 , fm ⁴	C_{12} , fm ¹⁰

Why: The Lennard-Jones row is worth pausing on: `LennardJones` takes C_6 *first*, even though the potential is universally called “12-6”. Reading the two the other way round silently swaps the two coefficients, and since both are positive nothing downstream complains — the tuning still converges, onto the wrong labels.

case	potential	p_1 obtained	p_1 published	p_2 obtained	p_2 published	deviation
nn	well	1.109228	1.1096	0.391087	0.3918	0.18%
nn	Pöschl-Teller	0.907086	0.9071	0.799622	0.7991	0.07%
nn	Gaussian	1.211903	1.2121	0.567906	0.5672	0.12%
nn	Lennard-Jones	9.760657	9.8667	3.005560	3.0884	2.68%
unit.	well	1.233701	1.2337	1.000000	1.0000	0.00%
unit.	Pöschl-Teller	1.000000	1.0000	2.000000	2.0000	0.00%
unit.	Gaussian	1.342002	1.3420	1.435224	1.4349	0.02%
unit.	Lennard-Jones	0.264625	0.2646	0.000341	0.00034	0.00%
deut.	well	1.758644	1.7575	0.499215	0.5000	0.16%
deut.	Pöschl-Teller	1.423895	1.4388	0.886001	0.8631	2.65%
deut.	Gaussian	1.910904	1.9102	0.674817	0.6754	0.09%
deut.	Lennard-Jones	6.846862	6.8147	0.915045	0.9049	1.13%

The two rows in bold are the **D1** divergence: the published parameters achieve $r_0 = 2.71$ and 1.73 , while Table 2 of the same paper asks for 2.70 and 1.70 . Refining the grid does not move these deviations. The third row above 0.2%, (deut., Lennard-Jones) at 1.13%, is not of that kind: there the published parameters do reach the target, and the gap is the flat direction in (C_{12}, C_6) discussed above.

SYSTEMS — Table 1, the two real systems

system	a	r_0	measured E	inferred $\hbar^2/2m_r$
deuteron	5.4112 fm	1.7436 fm	−2.224 MeV	41.462 MeV·fm ²
⁴ He dimer	90.4 Å	8.0 Å	−1.62 mK	12.095 K·Å ²